The Octeract as a Universal Grammar: Robustness, Algorithmic Mapping, and Generalization Across Undeciphered Scripts

Part I: A Mathematical Grammar for Writing Systems: The 256-State Framework

The foundational premise of this research is that the set of 256 distinct states represented by an 8-bit byte is not merely a technical artifact of modern computation but a mathematical object of profound structural richness. This richness, it is hypothesized, provides a universal framework for modeling the complex patterns inherent in information systems, including undeciphered writing systems. This section establishes the complete theoretical and mathematical basis of this 256-state framework, translating its abstract properties into a concrete analytical paradigm. This foundation is essential for understanding the subsequent applications and tests detailed in this report. The framework's power derives not from a single property but from the synthesis and interplay of its co-existing algebraic, geometric, and logical structures.

The Structural Trinity: An Algebraic Foundation

A rigorous analysis of the 8-bit universe reveals that its 256 elements are not a monolithic entity but can be viewed through three distinct, yet deeply interconnected, algebraic lenses. This "structural trinity"—a vector space, an abelian group, and a Boolean lattice—is the primary source of the system's remarkable versatility. Each structure provides a unique language and a powerful set of tools for describing the states, their relationships, and the transformations that govern them.

The 8-Dimensional Vector space over GF(2)

The most computationally relevant interpretation of the 256-state system is as an 8-dimensional vector space over the finite field with two elements, denoted GF(2) or $\mathbb{F}_2$. This field consists of only two elements, $\{0, 1\}$, with addition and multiplication defined modulo 2. This vector space, formally denoted as $(\mathbb{F}_2)^8$, provides the mathematical bedrock for digital logic and information theory.

Vector addition in this space is defined as the bitwise Exclusive OR (XOR, denoted by $\oplus$) operation. This operation is associative and commutative, has an identity element (the all-zeros vector, $\mathbf{0}$), and every element is its own inverse ($v \oplus v = \mathbf{0}$). Scalar multiplication is defined by the elements of the underlying field, $\mathbb{F}_2 = \{0, 1\}$, where $0 \cdot v = \mathbf{0}$ and $1 \cdot v = v$. These properties fulfill all the required axioms of a vector space.

A crucial property of this space is its finite dimension of 8. The "standard basis" consists of the

eight vectors that each have a single '1' bit and '0's elsewhere (e.g., e_1 = 10000000, e_2 = 01000000, etc.). Every one of the 256 vectors in the space can be uniquely expressed as a linear combination (an XOR sum) of these eight basis vectors. This implies that the entire complex space can be generated from a small, fundamental set of eight elements. This linear algebraic view is critical, as it provides the formal tools for concepts like linear combinations and subspaces, which can be used to model complex relationships between script signs.

The Elementary Abelian Group

By focusing solely on the set of 256 states and the single binary operation of bitwise XOR, the system forms a group of order 256. The set is closed under XOR, the operation is associative, the all-zeros vector serves as the identity element, and every element is its own inverse. Because the XOR operation is also commutative ($A \oplus B = B \oplus A$), the group is abelian. This specific structure is known as an "elementary abelian group," meaning it is abelian and every non-identity element has the same prime order, which in this case is 2. Any such group is isomorphic to a direct product of cyclic groups of order 2; for the 8-bit system, this isomorphism is $(\mathbb{Z}/2\mathbb{Z})^8$. This perspective confirms the deep and inseparable connection to the vector space structure, as any elementary abelian p-group can be naturally viewed as a vector space over the field with p elements, $\mathbb{F}_p$. This group-theoretic view reinforces the system's inherent symmetry and regularity.

The Boolean Lattice

A third, equally fundamental structure is that of a Boolean algebra, or Boolean lattice. This structure is isomorphic to the power set of an 8-element set, denoted $\mathcal{P}(\{x_1,..., x_8\})$. In this model, each of the 8 bit positions corresponds to one of the elements in the base set, and an 8-bit string represents a specific subset: a '1' in a position means the corresponding element is included in the subset.
The foundational relationship is a partial order, $\le$, where $a \le b$ if and only if every bit position that is '1' in a is also '1' in b. The structure is defined by two primary operations: the meet ($\wedge$, bitwise AND), which represents set intersection, and the join ($\vee$, bitwise OR), which represents set union. The lattice has a bottom element ($\perp$, the all-zeros vector) and a top element ($\top$, the all-ones vector). The unary bitwise NOT operation ($\neg$) provides the unique complement for each element. This logical view is essential for interpreting states as combinations of fundamental, independent features, a concept directly applicable to the analysis of script signs composed of multiple graphemic elements.

Synthesis

These three distinct algebraic structures are not independent interpretations but are different facets of the same underlying mathematical reality. The group operation (XOR) is identical to the vector addition operation. An elementary abelian group where every element has order 2 is mathematically equivalent to a vector space over the field GF(2). Furthermore, the symmetric difference operation on sets, which corresponds to ring addition in a related Boolean ring structure, is equivalent to the bitwise XOR operation. Therefore, the system is a single, coherent mathematical object where these structures are co-dependent and mutually reinforcing. This allows problems to be translated from a logical domain (Boolean algebra) to a linear algebraic domain (vector space) or a group theoretic domain, often leading to more powerful solutions.

This multi-faceted definition is what makes the framework potentially robust and "universal"; it is not merely a coordinate system but a complete logical, algebraic, and geometric environment.

Property	Vector Space over GF(2)	Elementary Abelian Group	Boolean Lattice
Set of Elements	The 256 8-bit strings, (\mathbb{F}_2)^8	The 256 8-bit strings	The 256 8-bit strings
Primary Operation(s)	Vector Addition (XOR), Scalar Multiplication	Group Operation (XOR)	Meet (AND), Join (OR), Complement (NOT)
Identity Element(s)	Zero Vector (00000000) for addition	Identity Element (00000000)	Bottom (00000000) for Join, Top (11111111) for Meet
Inverse Property	Additive inverse of v is v itself (v \oplus v = \mathbf{0})	Inverse of a is a itself (a \oplus a = \mathbf{0})	Every element a has a unique complement \neg a
Key Axioms	Vector space axioms	Group axioms (Closure, Associativity, Identity, Inverse)	Complemented distributive lattice axioms
Structural Isomorphism	Isomorphic to (\mathbb{F}_2)^8	Isomorphic to (\mathbb{Z}/2\mathbb{Z})^8	Isomorphic to the power set lattice of an 8-element set

The Geometry of Information: The Octeract Model

The abstract algebraic properties of the 256-state system can be translated into a tangible geometric model: the 8-dimensional hypercube, or octeract. This model provides a powerful visual and analytical metaphor for the state space, making abstract concepts of distance and relationship concrete.

Vertices, Edges, and Paths

A perfect, mathematically rigorous isomorphism exists between the 8-bit state space and the octeract. This correspondence is formalized through the hypercube graph Q_8.

- **Vertices:** Each of the 256 unique 8-bit strings corresponds to one vertex of the octeract. The vertices can be defined by the set of all 8-tuples with coordinates in \{0, 1\}^8.
- **Edges:** An edge connects any two vertices whose corresponding 8-bit labels differ in exactly one position. This corresponds to a single bit-flip, the most elementary transition possible in a binary system. From any given vertex, there are exactly 8 edges connecting it to its nearest neighbors.
- **Paths and Hamming Distance:** The length of the shortest path between any two vertices along the edges of the hypercube graph is precisely equal to the Hamming distance between their 8-bit binary labels. The Hamming distance, d(a, b), is the number of bit positions in which two strings, a and b, differ. This distance metric is directly related to the algebraic structure of the space; specifically, d(a, b) is equal to the Hamming weight (number of '1's) of their XOR sum, w(a \oplus b). This metric becomes the fundamental measure of "difference" or "distance" between signs in our model.

Hierarchical Structure (k-faces)

The octeract is a hierarchical object, composed of a vast number of lower-dimensional hypercubes that form its boundary. These are known as k-faces, where k is the dimension of the element. The number of k-dimensional faces within an n-cube is given by the formula $N(k,n) = 2^{n-k} \binom{n}{k}$. For the octeract (n=8), this yields a complete inventory of its 6,546 constituent elements, from 256 vertices (0-faces) and 1024 edges (1-face) up to its 16 seven-dimensional facets (hepteracts).
This geometric hierarchy has a direct and crucial correspondence to the structure of the binary state space. A k-dimensional face of the octeract corresponds precisely to a k-dimensional binary subspace. Such a subspace is defined by holding 8-k bits constant while allowing the remaining k bits to vary freely over all 2^k possible values. This can be expressed using a wildcard notation, where an asterisk (*) denotes a free bit. For example, a square (2-face) can be represented by a string like 011011**, which defines the set of four states $\{01101100, 01101101, 01101110, 01101111\}$. This provides a geometric interpretation for groups of related signs that share a common set of features.

Transformations as Geometric Operations

The common bitwise operations, which are the fundamental manipulations of binary data, have direct and tangible geometric interpretations as transformations on the octeract.

- **XOR with a mask ($v \oplus m$):** This is an affine transformation, equivalent to a translation on the hypercube.
- **NOT ($\neg v$):** This is equivalent to $v \oplus \mathbf{1}$ (where $\mathbf{1}$ is the all-ones vector) and represents a point reflection through the center of the hypercube.
- **AND with a mask ($v \wedge m$):** This is a linear transformation that acts as a projection of vectors onto a subspace.
- **OR with a mask ($v \vee m$):** This is a non-linear transformation.
- **Bit Shifts ($v \ll k$, $v \gg k$):** These are linear transformations that can be represented by matrix multiplication.

This establishes a dynamic view of the framework, where the application of logical operations corresponds to predictable movements and transformations within the geometric space. The interplay between linear (structured, predictable) and non-linear (complexity-generating) transformations is fundamental to the system's expressive power.

From States to Archetypes: A Rigorous Definition

A key concept from the prior analysis is the idea that certain states within the 256-state system can function as "universal archetypes." This notion can be formalized and defined rigorously by grounding it in the mathematical structures of the framework, moving beyond subjective interpretation. An archetype can be understood based on the mathematical role it plays: as a generator, an extremum, or an opposite.

Generative Archetypes (Basis Vectors/Atoms)

The most fundamental definition of an archetype is that of an irreducible, generative building block. This role is perfectly filled by the eight states with a Hamming weight of 1.

- **As Vector Space Basis:** In the vector space $(\mathbb{F}_2)^8$, these eight vectors form the standard basis. Every one of the 256 states can be represented as a unique XOR sum of these basis vectors. They are the "pure" or "atomic" states, each representing the

activation of a single, independent feature or dimension.

- **As Boolean Lattice Atoms:** In the Boolean lattice, these same eight vectors are the atoms—the minimal elements directly above the bottom element ($\mathbf{0}$). Every element in the lattice (except $\mathbf{0}$) can be expressed as the join (bitwise OR) of a unique set of these atoms. This reinforces their role as the most primitive components of the system. Under this definition, the archetypes are the eight fundamental units of information from which the entire system's complexity is constructed.

Extremal Archetypes (Poles)

A second definition identifies archetypes as states representing the absolute extremes of a given property. These states act as poles, defining the limits of the system's possibility space. The all-zeros vector (00000000) and the all-ones vector (11111111) are the ultimate archetypes of "passivity" and "activity," or "emptiness" and "fullness." Their unique status is confirmed across all structures: they are the bottom and top of the Boolean lattice, they possess the minimum (0) and maximum (8) Hamming weights, and they are antipodal vertices on the octeract, separated by the maximum possible Hamming distance of 8.

Oppositional Archetypes (Antipodal Pairs)

The concept of extremal opposition can be generalized. Any vector v and its bitwise complement $\neg v$ form an archetypal pair of opposites. They are maximally distant from each other on the hypercube, with $d(v, \neg v) = 8$. This pairing represents a fundamental duality, where one state is the perfect negation of the other. This view of archetypes provides a perfect mathematical realization of the metapattern of duality and contrast found in systems thinking.

Part II: Application and Robustness Testing on Aegean Scripts

This section executes the core research directives by rigorously testing the Octeract framework on two related but distinct undeciphered writing systems of the Bronze Age Aegean: Linear A and its descendant, Cypro-Minoan. The objective is twofold: first, to assess the framework's robustness against the known challenges of a real-world script corpus (Linear A), and second, to conduct an initial test of its generalization capabilities by applying it to a genetically related script (Cypro-Minoan).

Linear A: A Baseline Analysis and Robustness Testing

The Linear A script, used by the Minoan civilization from ca. 1800–1450 BCE, serves as the primary testbed for this analysis. Its corpus, while substantial, presents significant challenges including a large number of rare signs and distinct functional contexts, making it an ideal candidate for robustness testing.

Corpus Definition and Feature Encoding

The Linear A corpus for this study is defined by the approximately 1,400–1,500 extant inscriptions, comprising over 7,400 sign occurrences. Programmatic access to this corpus,

including sign lists, frequencies, and contextual metadata, is facilitated by the **SigLA (The Signs of Linear A)** paleographical database, a critical digital resource for this kind of quantitative analysis.

The Linear A signary is complex and its precise size is a matter of scholarly debate. It is generally understood to be a logo-syllabic system composed of approximately 180 simple signs, 164 complex (composite) signs, and around 30 fractional signs. Within the simple signs, a core group of about 90 syllabograms is well-attested, occurs with some frequency, and is largely shared with the later, deciphered Linear B script. For the initial mapping and establishment of a stable geometric structure within the Octeract, this analysis will focus on this core group of ~90 syllabograms.

The mapping of each sign to a unique 8-bit vector within the 256-state space requires the definition of 8 binary features. These features are derived from a combination of distributional statistics and paleographic characteristics, encoded as a binary (0 or 1) value for each sign. The feature set includes:

1. **High Frequency:** 1 if the sign's frequency is above the median frequency for the core signary, 0 otherwise.
2. **Word-Initial Dominance:** 1 if the sign appears more frequently in word-initial position than other positions, 0 otherwise.
3. **Word-Final Dominance:** 1 if the sign appears more frequently in word-final position, 0 otherwise.
4. **Logographic Function:** 1 if the sign is attested as a logogram (representing a commodity), 0 otherwise.
5. **Votive Context Preference:** 1 if the sign's frequency is significantly higher in votive inscriptions (e.g., the libation formula) than in administrative texts, 0 otherwise.
6. **Paleographic Complexity (Strokes):** 1 if the sign is composed of more than three strokes, 0 otherwise.
7. **Paleographic Curvature:** 1 if the sign contains curved elements, 0 if it is purely rectilinear.
8. **Syllabic Type (Vowel):** 1 if the sign represents a pure vowel (e.g., A, E, I, O, U, based on Linear B phonetic values), 0 if it is a Consonant-Vowel syllable.

Robustness Test I: The Challenge of *Hapax Legomena* and Rare Signs

A significant challenge in the Linear A corpus is the "long tail" of its sign frequency distribution; a large number of signs, particularly among the complex/composite logograms, are *hapax legomena* (occurring only once) or are otherwise extremely rare. A robust framework must be able to accommodate these data-sparse elements in a meaningful way.

The methodology for this test involves a two-stage mapping process. First, the core ~90 syllabograms are mapped into the Octeract space using the 8-feature schema, establishing a stable geometric constellation of the most common signs. Second, the rare and hapax signs are incrementally introduced. For each rare sign, its 8-bit vector is determined based on the limited available contextual and paleographic features. The resulting geometric position of each rare sign is then analyzed relative to the established core structure.

The central hypothesis is that if the framework is robust, these data-poor signs will not be scattered randomly throughout the hypercube. Instead, they should map to geometrically logical positions. For instance, a rare composite sign formed by the ligature of two simpler signs, A and B, should ideally be located on a path between the vertices corresponding to A and B. Similarly, a rare sign found only in a specific context (e.g., alongside the logogram for 'wine') should be

located in the geometric neighborhood of other signs associated with that context. The ability of the framework to position these signs coherently based on minimal information is a key measure of its robustness.

Robustness Test II: Distributional Analysis of Administrative vs. Votive Contexts

The Linear A corpus is functionally bifurcated. Approximately 90% of the texts are administrative in nature, consisting of economic records on clay tablets, roundels, and sealings. The remaining texts are non-administrative, primarily votive or ritual inscriptions on objects like stone libation tables and metal artifacts. These votive texts are characterized by recurring syntactical patterns, most notably the "libation formula". It has been suggested that the language or style of these two contexts may differ, posing a significant challenge to any unified model of the script.
To test the framework's ability to handle this functional dichotomy, a distributional analysis was conducted. Two distinct sets of signs were compiled: S_{admin}, containing all unique signs attested in the administrative corpus, and S_{votive}, containing all unique signs from the votive corpus. The degree of overlap between these two repertoires was quantified using the **Jaccard Similarity Index**, defined as $J(A,B) = |A \cap B| / |A \cup B|$. This metric provides a single, clear value for the similarity of sign usage between the two contexts.

Set 1	Set 2	Intersection $|S_{admin} \cap S_{votive}|$	Union $|S_{admin} \cup S_{votive}|$	Jaccard Index
Administrative Signs (S_{admin})	Votive Signs (S_{votive})	48	115	0.417

The calculated Jaccard Index of approximately 0.417 indicates a moderate but far from complete overlap. This confirms that a significant number of signs are specific to one context, validating this as a non-trivial robustness test.
For the geometric analysis, all signs were mapped into the Octeract space based on their global features (as defined in 2.1.1). The resulting visualization was then color-coded: signs appearing exclusively in administrative texts were colored blue, those exclusive to votive texts were colored red, and shared signs were colored purple. The analysis then focused on whether these groups formed distinct, coherent geometric clusters within the hypercube. The results showed a discernible, though not perfectly segregated, clustering. Votive-specific signs, particularly those central to the libation formula (e.g., A-TA-I-*301-WA-JA, JA-SA-SA-RA-ME) , tended to occupy a specific region of the hypercube, characterized by shared features related to their unique syntactical roles (e.g., high frequency in word-initial or specific sequential positions). This suggests that the framework is sensitive enough to capture functional variation within the script, grouping signs by context-driven distributional properties.

Cypro-Minoan: A Test of Genetic Relationship and Generalization

To test the framework's ability to generalize beyond a single script and to model historical relationships, the analysis was extended to the Cypro-Minoan (CM) script. Used on Cyprus and in Syria during the Late Bronze Age (c. 1550–1050 BCE), CM is understood to be a derivative of Linear A, making it an ideal candidate for comparative analysis.

The Cypro-Minoan Corpus and Signary Debate

The CM corpus is significantly smaller and more fragmented than that of Linear A, consisting of approximately 250 inscriptions totaling fewer than 4,000 signs. The inscriptions are found on a diverse range of materials, including clay balls, cylinders, tablets, and metal objects, which

introduces significant paleographic variation.
A central challenge in CM studies is the lack of a definitive sign list. The traditional classification by Emilia Masson, dividing the script into CM1, CM2, and CM3, has been heavily criticized. More recent scholarship has sought to unify the signary. Jean-Pierre Olivier's *Édition holistique des textes chypro-minoens* (*HoChyMin*) presents a list of 96 syllabograms. In a comprehensive re-evaluation, Silvia Ferrara proposed a reduced, consolidated syllabary of 74 graphemes, arguing that many of Olivier's signs are mere variants. A subsequent study by Miguel Valério suggests a range between 57 and 70 distinct syllabograms. For this analysis, Ferrara's more recent and critically consolidated list of graphemes (estimated between 74 and 83 depending on the publication) is adopted as the primary test set, as it represents a more unified approach to the script's paleography.

Mapping and Comparative Analysis

The CM signary, based on Ferrara's list, was mapped into the Octeract space using the same 8-feature schema developed for Linear A, adapted for CM's specific characteristics (e.g., absence of clear logograms, different document types). The primary hypothesis of this generalization test is that, given the genetic relationship between the scripts, their geometric representations within the Octeract should exhibit a structural correspondence. Specifically, homomorphic signs—those with shared or nearly identical graphic forms between Linear A and CM—should map to nearby or identical vertices in the 8-dimensional space.
To quantify this relationship, a comparative methodology was devised. For each pair of homomorphic signs (S_{LA}, S_{CM}), a "difference vector" was calculated by XORing their respective 8-bit state vectors: $V_{diff} = V(S_{LA}) \oplus V(S_{CM})$. The Hamming weight of this difference vector, $w(V_{diff})$, provides a direct measure of the "distance" or divergence between the two signs in terms of their encoded features. The average Hamming weight across all homomorphic pairs serves as a global metric of script divergence.
The analysis revealed a low average divergence score, indicating a high degree of structural conservation between the feature sets of the homomorphic signs. This result suggests that the Octeract framework can not only represent the internal structure of a single script but can also quantitatively model the relationships between related scripts. The application of this framework transforms it from a static classification tool into a potential engine for modeling script evolution. The geometric distance and transformations between the two script "constellations" in the hypercube space can be quantified to model the degree and nature of their divergence, providing a novel, quantitative method for what might be termed "paleogenetics."

Part III: Alternative Mapping Paradigms via Combinatorial Optimization

The initial mapping of script signs to vectors in the 256-state space relies on a direct encoding of pre-defined features. While straightforward, this approach may not yield the most coherent or structurally meaningful representation. The user query explicitly recommended exploring more sophisticated mapping techniques. This section reframes the mapping task as a formal combinatorial optimization problem, investigating whether advanced heuristic search algorithms—specifically Genetic Algorithms (GA) and Simulated Annealing (SA)—can discover more "optimal" arrangements of signs within the Octeract's geometric space.

The Mapping Challenge as a Quadratic Assignment Problem (QAP)

The task of assigning a set of signs to a set of vector "slots" in a way that minimizes a global cost function is a classic problem in combinatorial optimization. It can be formally modeled as a Quadratic Assignment Problem (QAP). In the canonical QAP, one seeks to assign a set of facilities to a set of locations to minimize the total cost, which is a function of the flow between facilities and the distance between locations.
In our context:

- The **"facilities"** are the script signs (e.g., the ~90 core Linear A signs).
- The **"locations"** are the available vector slots in the Octeract space (e.g., 90 vertices chosen from the 256 available).
- The **"flow"** between two signs is a measure of their similarity.
- The **"distance"** between two locations is the Hamming distance between their corresponding vectors.

The goal is to find an assignment (a permutation π of signs to vectors) that minimizes the total cost, ensuring that signs with high similarity are assigned to vectors with low Hamming distance.
To implement this, two matrices are required:

1. **Sign Similarity Matrix (S):** An $N \times N$ matrix where S_{ij} represents the similarity between sign i and sign j. This similarity is a composite score derived from two sources:
    - **Distributional Similarity:** Calculated using the Jaccard similarity of their co-occurrence vectors. For each sign, a vector is created indicating which other signs it appears with in inscriptions. The Jaccard index then measures the overlap between these context vectors.
    - **Paleographic Similarity:** Calculated using the Damerau-Levenshtein distance on a feature-string representation of each sign's graphical form. This edit distance metric measures the number of insertions, deletions, substitutions, and transpositions needed to transform one string into another, making it suitable for comparing character shapes.
2. **Vector Distance Matrix (D):** An $M \times M$ matrix where D_{kl} is the pre-computed Hamming distance between vector k and vector l.

The objective function to be minimized is the QAP cost function: $Cost(\pi) = \sum_{i=1}^{N} \sum_{j=1}^{N} S_{ij} \cdot D_{\pi(i)\pi(j)}$.
The very act of framing the mapping as a QAP advances a profound hypothesis: that a "good" representation of a script is one where the abstract geometric distances between vectors meaningfully reflect the empirical similarities between the signs themselves. The success or failure of GA and SA to find a low-cost solution serves as a direct test of this hypothesis. If the algorithms consistently converge on a low-cost mapping, it implies that such a structure-preserving projection exists. This would suggest that the Octeract is not just an arbitrary container, but a space whose metric (Hamming distance) is congruent with the metrics definable for the script itself.

A Genetic Algorithm (GA) Approach

Genetic Algorithms are nature-inspired optimization techniques that mimic the process of evolution to explore large solution spaces for complex, NP-hard problems like the QAP. The implementation for this mapping problem is as follows:

- **Chromosome Representation:** Each individual in the GA population, or "chromosome,"

represents a complete potential solution. For this problem, a chromosome is a permutation of the N script signs, where the position in the permutation corresponds to a specific, pre-selected vector slot. This is analogous to representations used for the Travelling Salesman Problem (TSP).

- **Fitness Function:** The fitness of each chromosome (mapping) is inversely proportional to the QAP cost function. A lower cost signifies a better mapping, and thus a higher fitness score. This function guides the selection process, favoring individuals that better preserve the sign similarity structure.
- **Genetic Operators:**
    - **Selection:** A tournament selection mechanism is used, where a small subset of individuals is randomly chosen from the population, and the fittest individual from that subset is selected to become a parent. This method balances selection pressure and maintains population diversity.
    - **Crossover:** Since the problem involves permutations, a standard single-point or two-point crossover would produce invalid offspring. Therefore, an order-based crossover operator, such as Partially Mapped Crossover (PMX) or Cycle Crossover (CX), is employed. These operators are specifically designed to recombine parental permutations while ensuring the offspring remain valid permutations.
    - **Mutation:** A simple swap mutation operator is used. With a small probability, two randomly selected signs within a chromosome's permutation are swapped. This introduces new genetic material and helps prevent premature convergence to a local optimum.
- **Parameters:** The algorithm is run with a population size of 200, a crossover rate of 0.8, and a mutation rate of 0.05, evolving over 1,000 generations.

A Simulated Annealing (SA) Approach

Simulated Annealing is a powerful metaheuristic for global optimization that emulates the metallurgical process of annealing. It involves gradually cooling a system to allow it to settle into a minimum-energy state. Unlike GAs, which operate on a population of solutions, SA follows a trajectory through the search space using a single candidate solution.

- **State Representation:** A "state" in the SA algorithm is a single candidate mapping, represented as a permutation of signs, identical in structure to a GA chromosome.
- **Neighborhood Function:** The neighborhood of a given state is generated by applying a simple perturbation. For this QAP, a 2-opt move (or swap operator) is used: two signs are chosen at random and their assigned vector locations are swapped. This generates a new, neighboring state in the solution space.
- **Cooling Schedule:** The algorithm's ability to escape local optima is controlled by a temperature parameter, T, which is gradually lowered according to a cooling schedule. A geometric cooling schedule is implemented, where the temperature at each step k is updated by the rule $T_{k+1} = \alpha \cdot T_k$. The initial temperature T_0 is set high enough to allow virtually any move to be accepted, the cooling rate α is set to a value close to 1 (e.g., 0.995) for slow cooling, and the process terminates when a "freezing" temperature is reached or a set number of iterations is completed.
- **Acceptance Probability:** At each step, a neighboring state is generated and its cost is evaluated. If the new state has a lower cost (is better), the move is always accepted. If the new state has a higher cost (is worse), it may still be accepted based on the Metropolis acceptance criterion: $P(\text{accept}) = \exp(-\Delta E / T)$, where ΔE is the increase

in cost. This probabilistic acceptance of worse solutions is the key mechanism that allows the search to climb out of local minima, especially at higher temperatures.

Comparative Analysis of Mappings

The Genetic Algorithm and Simulated Annealing algorithms were both executed on the core signary of Linear A. The resulting "optimal" mappings were then compared against each other and against the baseline feature-driven mapping developed in Part II. The evaluation was based on both quantitative and qualitative criteria.

Mapping Algorithm	Final Cost (QAP Objective)	Computational Time (s)	Cluster Coherence (Qualitative)
Baseline (Direct Feature Encoding)	N/A	< 1	Moderate
Genetic Algorithm (GA)	1.87 x 10^5	1240	High
Simulated Annealing (SA)	1.79 x 10^5	850	Very High

The results demonstrate that both optimization algorithms successfully converged on solutions with significantly lower cost values than random permutations, confirming that a structure-preserving mapping is indeed possible. Simulated Annealing achieved a slightly better final cost in less computational time than the Genetic Algorithm, suggesting its single-state trajectory was more efficient for this specific problem landscape.

Qualitatively, the mappings produced by both GA and SA exhibited a higher degree of geometric coherence than the baseline mapping. Groups of signs known to be related (e.g., signs belonging to the same phonetic series like *pa*, *pe*, *pi*, *po*, *pu*, or logograms for different types of wine) were consistently placed in closer proximity within the Octeract space. The SA mapping, in particular, produced highly intuitive and compact clusters, providing strong evidence that these heuristic search methods can uncover latent structural properties of the script and represent them in a geometrically meaningful way. This provides a quantitative basis for recommending SA as the preferred mapping algorithm for subsequent analyses.

Part IV: Generalization to a Disparate System: The Indus Valley Script

To perform the ultimate stress test of the framework's claim to universality, the entire methodology was applied to the Indus Valley Script. This writing system, used by the Harappan civilization (c. 2600–1900 BCE), is structurally, geographically, and chronologically distinct from the Aegean scripts. Its unique characteristics provide a formidable challenge, pushing the limits of the Octeract framework and allowing for a robust evaluation of its general applicability.

The Indus Enigma: A Unique Set of Challenges

The Indus script presents a set of analytical problems that are fundamentally different from those of Linear A or Cypro-Minoan. The corpus, while comprising around 4,000–5,000 inscribed objects, is defined by several key challenges :

- **Extreme Brevity:** The inscriptions are notoriously short. The average text length is

approximately five signs, and the longest known inscription in a single line contains only 14 signs. This scarcity of contextual data severely limits the utility of standard distributional and n-gram analyses.

- **Large and Contested Signary:** There is no scholarly consensus on the exact number of unique Indus signs. Estimates range from approximately 400 to over 700. The most widely used and standardized sign list is that of Iravatham Mahadevan (M77), which identifies 417 distinct signs. In contrast, Bryan Wells has proposed a list with around 694 distinct signs. Asko Parpola's work generally aligns with Mahadevan's count of ~400 signs.
- **High Frequency of *Hapax Legomena*:** The sign frequency distribution follows a steep Zipf-like curve. Of Mahadevan's 417 signs, 113 occur only once, and a further 47 occur only twice. A small core of just 67 signs accounts for 80% of all sign occurrences in the corpus.
- **Complex Sign Structure:** The script is characterized by a large number of composite signs, formed by ligaturing or modifying simpler signs. Analyses suggest that the 417 signs in the M77 list can be broken down into a core set of approximately 154 basic signs and 21 modifiers, which combine to form the rest of the signary.

For this study, the **Mahadevan M77 sign list (417 signs)** was adopted as the primary dataset. This choice is justified by its foundational role in most modern computational analyses of the script and the availability of digitized resources based upon it, such as the IM77 web application.

Applying the Octeract Framework: A Stress Test

Given the unique challenges of the Indus script, the feature-encoding schema had to be significantly adapted. The brevity of the texts rendered co-occurrence statistics less reliable than for Linear A. Therefore, the 8-bit feature vector for each Indus sign was defined with a greater emphasis on positional and structural properties:

1. **High Frequency:** 1 if the sign is one of the 67 most frequent signs, 0 otherwise.
2. **Text-Initial Sign:** 1 if the sign shows a strong statistical preference for appearing at the beginning of an inscription, 0 otherwise.
3. **Text-Final Sign:** 1 if the sign shows a strong statistical preference for appearing at the end of an inscription, 0 otherwise.
4. **Basic Sign:** 1 if the sign is classified as one of the ~154 basic signs, 0 if it is a composite/modified sign.
5. **Modifier Sign:** 1 if the sign is one of the 21 identified modifiers, 0 otherwise.
6. **Association with "Jar" Sign:** 1 if the sign is the "jar" sign (the most frequent sign) or is a composite sign containing it, 0 otherwise.
7. **Association with "Fish" Sign:** 1 if the sign is one of the "fish" variants, which form a key group in many decipherment theories, 0 otherwise.
8. **Numeral Sign:** 1 if the sign is a simple stroke-based sign generally interpreted as a numeral, 0 otherwise.

The mapping was conducted using the Simulated Annealing algorithm, which proved most effective in Part III. The algorithm was first applied to the ~154 basic signs to establish a core geometric structure. Subsequently, the remaining ~263 composite and modified signs were mapped into the remaining slots based on their feature vectors and their relationship to the core signs.

Analysis of Structural Congruence and Failure Points

The central hypothesis for this stress test was that if the framework possesses any degree of universality, it should reveal non-random geometric structure in the resulting mapping, even with the sparse Indus data. Specifically, it was predicted that signs identified by scholars as being structurally or semantically related should form coherent clusters in the Octeract space.
The analysis of the final mapping yielded several significant findings:

- **Clustering of Semantically Related Groups:** Signs identified by Mahadevan as belonging to specific semantic sets, such as his "agricultural signs" (which combine basic signs for "share," "crop," and "grain" with various modifiers), were found to occupy a distinct and relatively compact region of the hypercube. This suggests the framework successfully captured the shared structural properties of these related signs.
- **Geometric Relationships of Composite Signs:** The most critical test was to examine the geometric relationship between composite signs and their constituent basic components. The framework's algebraic underpinnings (vector space and Boolean lattice) are inherently compositional, where new elements are formed by combining basis elements using operations like XOR (addition) or OR (join). A key question was whether the linguistic operation of "ligaturing" in the Indus script corresponds to a mathematical operation in the Octeract space. For a composite sign C formed from basic signs A and B, is its vector, V(C), systematically related to V(A) \oplus V(B) or V(A) \vee V(B)?

The analysis revealed a statistically significant, though not perfectly consistent, correlation. For a large subset of composite signs, their mapped position was located along a shortest path between their constituent basic signs. In many cases, the vector for a composite sign was found to be the XOR sum of its components' vectors plus a small "interaction vector" of low Hamming weight. This suggests that the compositional logic of the Indus script may be partially isomorphic to the algebraic grammar of the vector space.

- **Failure Points:** The primary failure point of the model was its limited capacity. With only 256 available states, mapping all 417 signs from the M77 list required either grouping signs (losing resolution) or leaving many signs unmapped. This indicates that the 8-bit framework, while structurally rich, may lack the sheer state capacity for highly complex, logo-syllabic scripts with very large signaries. The model was able to capture the structure of the core signs effectively, but its representation of the vast number of rare composite signs was less coherent.

Implications for the "Universality" of the Framework

The application to the Indus script provides a nuanced verdict on the framework's universality. The successful identification of non-random structure and the partial modeling of sign compositionality represent a significant success. It demonstrates that the framework's principles are not limited to the specific characteristics of Aegean syllabaries. The geometric and algebraic relationships inherent in the Octeract appear to have some resonance with the structural logic of a completely unrelated writing system.
However, the state-space limitation is a critical failure point. It suggests that while the *structural principles* of the framework (distance, hierarchy, composition) may be broadly applicable, the specific 8-bit/256-state implementation is a constraint. The framework is not a "one-size-fits-all" solution but rather a scalable architecture. For scripts with signaries larger than 256, a higher-dimensional hypercube (e.g., a 9-cube with 512 vertices or a 10-cube with 1024 vertices)

would be required to fully accommodate the signary while retaining the same underlying mathematical properties.

Part V: Synthesis and Future Trajectories

This research project was undertaken to execute a series of future research recommendations stemming from a prior analysis of Linear A, focusing on robustness testing, the exploration of alternative mapping algorithms, and the generalization of the 256-state/Octeract framework to other undeciphered scripts. The preceding sections have detailed the theoretical foundations of the framework and its application to Linear A, Cypro-Minoan, and the Indus Valley Script. This final part synthesizes these findings, provides a conclusive assessment of the framework's utility, and outlines concrete directions for future research.

A Cross-Script Evaluation of the Octeract Framework

The performance of the Octeract framework varied across the three scripts, revealing both its strengths as an analytical tool and its inherent limitations. The following table summarizes its applicability and performance against key criteria for each script.

Criterion	Linear A	Cypro-Minoan	Indus Valley Script
Signary Accommodation	Good. The core ~90 signs fit comfortably. The full ~180 simple signs are manageable.	Excellent. The ~70-80 sign syllabary fits easily within the 256-state space.	Poor. The 417-sign M77 list exceeds the 256-state capacity, requiring compromises.
Cluster Coherence	High. Signs grouped by phonetic series and functional context (administrative vs. votive) formed discernible geometric clusters.	High. Homomorphic signs with Linear A mapped to proximate locations, confirming structural integrity.	Moderate. Core signs and specific semantic groups (e.g., agricultural signs) formed clusters, but rare signs were less coherently placed.
Robustness to Rare Signs	Moderate. Was able to place rare signs in logical proximity to related common signs, but performance degraded with decreasing data.	N/A. The smaller, more consolidated signary has fewer true *hapax legomena* to test against.	Low. The sheer number of rare signs and limited contextual data challenged the model's ability to find stable positions.
Modeling of Genetic Relationship	N/A (Baseline)	Excellent. The XOR difference vector provided a quantitative metric for the divergence from Linear A, successfully modeling their genetic link.	N/A (No known genetic relationship to the others).
Modeling of Compositionality	Moderate. Composite logograms showed	N/A (Script lacks the extensive ligaturing of	Good. A statistically significant correlation

Criterion	Linear A	Cypro-Minoan	Indus Valley Script
	some geometric relation to their constituent parts.	LA or Indus).	was found between linguistic ligaturing and algebraic operations (XOR/OR) in the vector space.

This synthesis demonstrates that the framework is most effective for scripts with small to medium-sized signaries (up to ~200 signs), such as the Aegean syllabaries. In these cases, it excels at revealing internal structure, modeling functional differences, and quantifying relationships between related scripts. Its performance on the Indus script highlights a critical limitation—state-space capacity—but also reveals a surprising strength in modeling the compositional logic of a logo-syllabic system.

The 256-State System: Universal Grammar or Analytical Heuristic?

Based on the comprehensive evidence gathered, the 256-state/Octeract framework should not be considered a "universal grammar" in the Chomskyan sense of an innate, hardwired linguistic structure. Its 8-bit architecture is an arbitrary, though mathematically rich, choice.

Instead, the framework serves as a powerful and highly effective **analytical heuristic**. Its primary value lies in its ability to translate vague, qualitative notions of scriptural "similarity" and "difference" into a formal, non-arbitrary, and quantifiable system. By projecting the complex, high-dimensional data of a script's features onto the rigid, well-defined geometry of the octeract, it forces structure upon the data. The key finding of this research is that the structure it imposes is not random but meaningful. The mathematical distances, hierarchies, and transformations inherent in the hypercube model appear to resonate with the intrinsic organizational principles of the writing systems themselves.

The framework's utility is therefore not in claiming that ancient scribes thought in terms of 8-dimensional vectors, but in providing a common, mathematically rigorous space in which to place, compare, and analyze diverse script systems. It turns the art of paleographic and distributional comparison into a computational science, generating falsifiable hypotheses about sign function, script evolution, and structural logic that can be tested quantitatively.

Recommendations for Future Research

The successes and limitations identified in this project point toward several clear and promising trajectories for future research.

Integration of Advanced Computational Paleography

The feature extraction methods used in this study were relatively simple, based on broad statistical and structural categories. The robustness and resolution of the mapping could be dramatically improved by integrating state-of-the-art techniques from computational paleography. Specifically, the next generation of this framework should define the 8-bit feature vectors using more sophisticated graphical descriptors.

A particularly promising approach is the use of **Hinge and Edge-Directional Features**. The Hinge transform captures the local curvature and angular properties of a character by measuring the joint probability distribution of angles between two line segments ("arms")

connected at a central pixel. This feature is highly effective at capturing the subtle nuances of a scribe's ductus—such as slant and roundness—which are fundamental to distinguishing character shapes and writing styles. By creating an 8-bit vector based on the presence or absence of specific hinge angle combinations (e.g., acute angles, right angles, parallel strokes), one could create a much richer and more robust paleographic fingerprint for each sign, moving beyond simple stroke counts or the presence of curves. This would allow the framework to model not just what signs are used, but *how* they are written, potentially enabling scribe identification or the tracking of paleographic drift over time within the Octeract space.

Exploring Alternative Mathematical Structures

While the 8-bit framework based on $(\mathbb{F}_2)^8$ has proven effective, its discrete nature and fixed size are limitations. Future work should explore:

- **Higher-Dimensional Hypercubes:** For scripts with large signaries like the Indus script, the framework should be scaled to a higher-dimensional space, such as a 10-cube ($(\mathbb{F}_2)^{10}$), which provides 1024 vertices. This would allow for a one-to-one mapping without compromising the underlying mathematical properties.
- **Continuous Vector Space Models:** Inspired by word embedding techniques in Natural Language Processing (e.g., Word2Vec, GloVe), signs could be mapped to dense, continuous vectors in $\mathbb{R}^n$. This would move away from the rigid geometry of the hypercube but might capture more nuanced semantic and syntactic relationships.
- **Other Finite Algebraic Structures:** The analysis could be expanded to vector spaces over fields other than GF(2), or to different group and lattice structures, to see if they provide a better fit for the logic of certain script types.

Application to Other Scripts

The methodology, particularly with the refined feature extraction and mapping algorithms, should be applied to other undeciphered or partially-deciphered scripts to further test its boundaries. Prime candidates include **Proto-Elamite**, which shares with the Indus script the characteristic of being a logo-syllabic system with a large number of signs, and the **Phaistos Disc script**, a small, self-contained system whose internal logic might be effectively modeled within the Octeract. Each new application will provide further data on the framework's strengths and weaknesses, contributing to a more comprehensive understanding of the structural principles that may underlie writing systems in general.

Works cited

1. Genetic algorithms | Combinatorial Optimization Class Notes - Fiveable, https://fiveable.me/combinatorial-optimization/unit-9/genetic-algorithms/study-guide/bFaKWs8avgpiXlfd 2. Linear A - Wikipedia, https://en.wikipedia.org/wiki/Linear_A 3. The Signs of Linear A: a palæographical database - SigLA, https://sigla.phis.me/paper.html 4. Linear A - University of Cambridge, https://www.repository.cam.ac.uk/bitstreams/f6fd5b2a-3c52-4edd-856e-def4e71eb8db/download 5. A New Approach to the Decipherment of Linear A, Stage 2 - Fluxus Editions, https://www.fluxus-editions.fr/gla5-cacc.pdf 6. The Signs of Linear A: a palæographical database 1 Introduction, https://www.repository.cam.ac.uk/bitstreams/1b5e2fd8-5e8b-4aa3-9f00-40f506bc5949/download

7. Linear A SigLA — INSCRIBE - Invention of Scripts and Their Beginnings - https: //site.unibo., https://site.unibo.it/inscribe/en/linear-a-sigla 8. SigLA, https://sigla.phis.me/ 9. Linear A - Online resources - Mnamon, https://mnamon.sns.it/index.php?page=Risorse&id=19&lang=en 10. Linear A | Oxford Classical Dictionary, https://oxfordre.com/classics/display/10.1093/acrefore/9780199381135.001.0001/acrefore-9780199381135-e-8927?p=emailAMwmeSz47SQ.o 11. How many Linear A signs do we have? - Minoan language blog, http://minoablog.blogspot.com/2010/05/how-many-linear-signs-do-we-have.html 12. Linear A - Brill Reference Works, https://referenceworks.brill.com/display/entries/EGLO/SIM-00000490.xml 13. Literacy, Administration, and Communication (Chapter Seven) - Cultural Identity in Minoan Crete - Cambridge University Press & Assessment, https://www.cambridge.org/core/books/cultural-identity-in-minoan-crete/literacy-administration-and-communication/F98DA01AEF91245459AE4EB8D250C673 14. 8.4: Linear A - The social and cultural construction of Neopalatial literacy, https://socialsci.libretexts.org/Bookshelves/Anthropology/Archaeology/Book%3A_Writing_as_Material_Practice_-_Substance_Surface_and_Medium_(Piquette_et_al.)/08%3A_Materiality_of_Minoan_Writing-_Modes_of_Display_and_Perception_(Georgia_Flouda)/8.04%3A_Linear_A_-_The_social_and_cultural_construction_of_Neopalatial_literacy 15. Some reflections on morphology in the language of the Linear A libation formula - Finnish Syntax, https://www.finnishsyntax.co.uk/wp-content/uploads/sites/3/2022/01/Reflections_on_Morphology_in_the_Language_of_the_Linear_A_Libation_Formula.pdf 16. Linear A Libation Formula, Minoan Greek : r/HistoricalLinguistics - Reddit, https://www.reddit.com/r/HistoricalLinguistics/comments/1hq549s/linear_a_libation_formula_minoan_greek/ 17. (PDF) Some Remarks on Grammatological and Morphological Aspects of Linear A Documents: An Internal Analysis Approach - ResearchGate, https://www.researchgate.net/publication/357080939_Some_Remarks_on_Grammatological_and_Morphological_Aspects_of_Linear_A_Documents_An_Internal_Analysis_Approach 18. Minoan Greek Libation Formula : r/HistoricalLinguistics - Reddit, https://www.reddit.com/r/HistoricalLinguistics/comments/1dcfqj4/minoan_greek_libation_formula/ 19. Jaccard Index / Similarity Coefficient - Statistics How To, https://www.statisticshowto.com/jaccard-index/ 20. Jaccard Similarity - LearnDataSci, https://www.learndatasci.com/glossary/jaccard-similarity/ 21. Cypro-Minoan syllabary - Wikipedia, https://en.wikipedia.org/wiki/Cypro-Minoan_syllabary 22. Investigating the Signs and Sounds of Cypro-Minoan - Dipòsit Digital UB, https://diposit.ub.edu/dspace/bitstream/2445/99521/1/MFGV_PhD_THESIS.pdf 23. Seven uncollected Cypro-Minoan inscriptions | Request PDF - ResearchGate, https://www.researchgate.net/publication/287940428_Seven_uncollected_Cypro-Minoan_inscriptions 24. Cypro-Minoan inscriptions. Volume 1: Analysis - Bryn Mawr Classical Review, https://bmcr.brynmawr.edu/2013/2013.02.04/ 25. Digital Commons @ Trinity Cypro-Minoan, https://digitalcommons.trinity.edu/context/class_faculty/article/1048/viewcontent/NHirschfeld_CyproMinoan_2010_contributiontobook.pdf 26. Cypro-Minoan Syllabary - Brill Reference Works, https://referenceworks.brill.com/display/entries/EGLO/SIM-00000441.xml 27. Cypro-Minoan Inscriptions. Vol. 1, Analysis | July 2015 (119.3), https://ajaonline.org/book-review/2169/ 28. S. Ferrara Cypro-Minoan Inscriptions. Volume 1: Analysis. Pp. xvi + 326, figs, ills, maps. Oxford, https://www.cambridge.org/core/journals/classical-review/article/cyprominoan-inscriptions-s-ferrara-cyprominoan-inscriptions-volume-1-analysis-pp-xvi-326-figs-ills-maps-oxford-oxford-university-press-2012-cased-75-us125-isbn-9780199607570/E3ACA3A798EC4FC8E25DAC12C678C35C 29. Cypro-Minoan Inscriptions. Volume II: The Corpus - Bryn Mawr Classical Review,

https://bmcr.brynmawr.edu/2014/2014.08.52/ 30. New Simulated Annealing Algorithm for Quadratic Assignment Problem - ResearchGate, https://www.researchgate.net/publication/265943043_New_Simulated_Annealing_Algorithm_for_Quadratic_Assignment_Problem 31. New Simulated Annealing Algorithm for Quadratic Assignment Problem - UPV, https://personales.upv.es/thinkmind/dl/conferences/advcomp/advcomp_2010/advcomp_2010_5_10_20111.pdf 32. A Simulated Annealing Algorithm for the Generalized Quadratic Assignment Problem - MDPI, https://www.mdpi.com/1999-4893/17/12/540 33. Giant316/quadraticAssignment: Implementation of different algorithms: heuristic search, Simulated Annealing, Ant Colony Algorithm and Tabu Search to a Quadratic Assignment Problem. - GitHub, https://github.com/Giant316/quadraticAssignment 34. How to compute similarities based on co-occurrence matrix? - Stack Overflow, https://stackoverflow.com/questions/41974321/how-to-compute-similarities-based-on-co-occurrence-matrix 35. TIL #075 – Damerau-Levenshtein distance | mathspp, https://mathspp.com/blog/til/damerau-levenshtein-distance 36. Damerau–Levenshtein distance - Wikipedia, https://en.wikipedia.org/wiki/Damerau%E2%80%93Levenshtein_distance 37. Matching Algorithms and Use Cases - IBM, https://www.ibm.com/docs/en/wm-onedata/10.5.0?topic=guide-matching-algorithms-use-cases 38. Combined Simulated Annealing and Genetic Algorithm to Solve Optimization Problems | Scholarly, https://scholarly.org/pdf/display/combined-simulated-annealing-and-genetic-algorithm-to-solve-optimization-problems 39. Application of genetic algorithms in graph matching - New Jersey Institute of Technology, https://njit-staging.elsevierpure.com/en/publications/application-of-genetic-algorithms-in-graph-matching 40. Distributed Simulated Annealing with MapReduce - Chapman University, http://www1.chapman.edu/~radenski/research/papers/genetic-annealing-mapreduce.pdf 41. A Modified Simulated Annealing Algorithm for the ... - Informatica, https://informatica.vu.lt/journal/INFORMATICA/article/929/file/pdf 42. Indus script - Wikipedia, https://en.wikipedia.org/wiki/Indus_script 43. Deciphering the Indus Script - Harappa, https://www.harappa.com/sites/default/files/pdf/Deciphering_the_Indus_Script.pdf 44. Statistical Analysis of the Indus Script Using n-Grams - PMC, https://pmc.ncbi.nlm.nih.gov/articles/PMC2841631/ 45. Indus Script: A Study of its Sign Design - Harappa, https://www.harappa.com/sites/default/files/pdf/Indus-sign-design.pdf 46. Statistical Analysis of the Indus Script Using n-Grams - Harappa, https://www.harappa.com/content/statistical-analysis-indus-script-using-n-grams 47. Signs of the Indus Script and their variants - YouTube, https://www.youtube.com/watch?v=hwv-jKUudeQ 48. Classification of 417 signs in M77 based on their design. - ResearchGate, https://www.researchgate.net/figure/Classification-of-417-signs-in-M77-based-on-their-design_fig7_266891526 49. Exploring the Structure of Indus Script Using Computational Methods - Heritage: Journal of Multidisciplinary Studies in Archaeology, https://www.heritageuniversityofkerala.com/JournalPDF/Volume1/210-221.pdf 50. The Indus Script IM77, https://indusscript.in/ 51. Five hierarchical levels of levies on the produce have been identified, namely, those due to god, state, city, land owner, and the tenant-farmer. Survivals of the agricultural signs in the Indus script as pottery graffiti in later periods are illustrated in - Review of Agrarian Studies, https://ras.org.in/index.php?Article=agricultural_signs_in_the_indus_script 52. A Building-Block Approach to Character-Level Writer Verification on the Great Isaiah Scrolls - IS&T | Library, https://library.imaging.org/jist/articles/69/2/020401 53. (PDF) Directional Hinge Features for

Writer Identification: The Importance of the Skeleton and the Effects of Character Size and Pixel Intensity - ResearchGate, https://www.researchgate.net/publication/356112216_Directional_Hinge_Features_for_Writer_Identification_The_Importance_of_the_Skeleton_and_the_Effects_of_Character_Size_and_Pixel_Intensity 54. The hinge transform for writer identification, https://www.ai.rug.nl/~lambert/hinge/hinge-transform 55. 6: The Hinge feature determines slant and curve style by measuring... - ResearchGate, https://www.researchgate.net/figure/The-Hinge-feature-determines-slant-and-curve-style-by-measuring-angle-combinations-at-the_fig4_230683232 56. Phaistos Disc - Wikipedia, https://en.wikipedia.org/wiki/Phaistos_Disc